

Module 3 Lab 2 Solution

1. Load the dataset `bike_day_raw.csv`, which has the regression target `cnt`. This dataset is hourly bike rentals in the citybike platform. The `cnt` column is the number of rentals, which we want to predict from date and weather data.
- Split the data into a training and a test set using `train_test_split`.
 - Use the `LinearRegression` class to learn a regression model on this data.
 - You can evaluate with the `score` method, which provides the R^2 or using the `mean_squared_error` function from `sklearn.metrics` (Challenge: You can also write it yourself in `numpy`).

```
In [3]: import pandas as pd
        from sklearn.model_selection import train_test_split

        data = pd.read_csv("bike_day_raw.csv")
        X = data.drop("cnt", axis=1)
        y = data.cnt

        display(data.head())

        X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
```

	season	mnth	holiday	weekday	workingday	weathersit	temp	atemp	hum	w
0	1	1	0	6	0	2	0.344167	0.363625	0.805833	
1	1	1	0	0	0	2	0.363478	0.353739	0.696087	
2	1	1	0	1	1	1	0.196364	0.189405	0.437273	
3	1	1	0	2	1	1	0.200000	0.212122	0.590435	
4	1	1	0	3	1	1	0.226957	0.229270	0.436957	

```
In [4]: from sklearn.linear_model import LinearRegression
import numpy

lr = LinearRegression().fit(X_train, y_train)

print(lr.score(X_train, y_train))
ybar = numpy.sum(y_train)/len(y_train)      # or sum(y)/len(y)
ssreg = numpy.sum((lr.predict(X_train)-ybar)**2) # or sum([ (yihat - ybar)*
sstot = numpy.sum((y_train - ybar)**2)      # or sum([ (yi - ybar)**2 for yi in
print(ssreg / sstot)

print(lr.score(X_test, y_test))

0.5328925529498699
0.5328925529498698
0.4991033756876271
```

1. Load the diabetes dataset using `sklearn.datasets.load_diabetes`.

- Scale the dataset (you can be creative and make pipelines)
- Apply LinearRegression,
- Apply Ridge and do grid search
- Apply Lasso and do grid search
- Visualize the coefficients.



```
In [5]: from sklearn.linear_model import Lasso, Ridge, LinearRegression
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.datasets import load_diabetes
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

diabetes = load_diabetes()

# create dataframe for easy boxplot
df = pd.DataFrame(diabetes.data, columns=diabetes.feature_names)
df.boxplot()

plt.figure()
plt.title("Target distribution")
plt.hist(diabetes.target, bins="auto")

X_train, X_test, y_train, y_test = train_test_split(diabetes.data,
                                                    diabetes.target)
```



```
In [6]: # With scaled data
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler().fit(X_train)
X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)


scores_lr = cross_val_score(LinearRegression(), X_train_scaled, y_train, cv=1)
print("Linear regression w/ scaling:", scores_lr.mean())
scores_ridge = cross_val_score(Ridge(), X_train_scaled, y_train, cv=10)
print("Ridge regression w/ scaling:", scores_ridge.mean())

from sklearn.model_selection import GridSearchCV
param_grid = {'alpha': np.logspace(-3, 3, 7)}
grid = GridSearchCV(Ridge(), param_grid, cv=10, return_train_score=True)
grid.fit(X_train_scaled, y_train)

res = pd.DataFrame(grid.cv_results_)
res.plot("param_alpha", ["mean_train_score", "mean_test_score"], logx=True)
plt.title("Ridge grid search")
```


Linear regression w/ scaling: 0.46894845696632537
 Ridge regression w/ scaling: 0.47052967021791836

Out[6]: Text(0.5, 1.0, 'Ridge grid search')




```
In [7]: lr = LinearRegression().fit(X_train, y_train)
plt.scatter(range(X_train.shape[1]), lr.coef_, c=np.sign(lr.coef_), cmap="bwr")
```

Out[7]: <matplotlib.collections.PathCollection at 0x1227edb20>



```
In [8]: ridge = grid.best_estimator_
plt.scatter(range(X_train.shape[1]), ridge.coef_, c=np.sign(ridge.coef_), cma
```

Out[8]: <matplotlib.collections.PathCollection at 0x1219131c0>



```
In [9]: param_grid = {'alpha': np.logspace(-3, 0, 13)}
print(param_grid)
```

```
{'alpha': array([0.001      , 0.00177828, 0.00316228, 0.00562341, 0.01      ,
                  0.01778279, 0.03162278, 0.05623413, 0.1      , 0.17782794,
                  0.31622777, 0.56234133, 1.          ])}
```

```
In [10]: grid = GridSearchCV(Lasso(normalize=True, max_iter=1e6), param_grid, cv=10, r
grid.fit(X_train, y_train)
```

/Users/gceran/opt/anaconda3/lib/python3.8/site-packages/sklearn/model_selection/_search.py:847: FutureWarning: The parameter 'iid' is deprecated in 0.22 and will be removed in 0.24.

```
warnings.warn(
```

```
Out[10]: GridSearchCV(cv=10, estimator=Lasso(max_iter=1000000.0, normalize=True),
                    iid=False,
                    param_grid={'alpha': array([0.001      , 0.00177828, 0.00316228, 0.
.00562341, 0.01      ,
                    0.01778279, 0.03162278, 0.05623413, 0.1      , 0.17782794,
                    0.31622777, 0.56234133, 1.          ])},
                    return_train_score=True)
```

```
In [11]: print(grid.best_params_)
print(grid.best_score_)
```


```
{'alpha': 0.1}
0.47667873696118457
```

```
In [12]: grid.score(X_test, y_test)
```

```
Out[12]: 0.43083075264606996
```

```
In [13]: lasso = grid.best_estimator_
plt.scatter(range(X_train.shape[1]), lasso.coef_, c=np.sign(lasso.coef_), cma
```

Out[13]: <matplotlib.collections.PathCollection at 0x1218f5b80>



In []: